

# STA9890: Regression

Jonathan Ng

```
library(tidyverse)
library(gridExtra)
library(reshape2)
library(stringr)
library(FSA)
library(ISLR)
library(glmnet)
library(randomForest)
library(kableExtra)
library(plotly)
library(rbenchmark)

set.seed(1)
data <- read_csv("OnlineNewsPopularity.csv") %>%
  select(-url, -timedelta) %>%
  sample_n(5000)
```

## 1. Mashable Online News Popularity

<https://archive.ics.uci.edu/ml/datasets/Online+News+Popularity>

### (a) Describe the response variable and the predictors. How was the data collected?

This dataset from *Mashable* summarizes sets of features about articles published on their website [www.mashable.com](http://www.mashable.com). The goal is to predict response variable, the number of shares in social networks for a given article, which is an indicator of popularity.

The data was collected solely from articles published on [www.mashable.com](http://www.mashable.com) for a two year period. The data does not contain the actual content of articles, but rather various summary statistics and metadata extracted from the published articles.

**The Response Variable: shares**

**The Predictor Variables:**

1. `n_tokens_title`: Number of words in the title
2. `n_tokens_content`: Number of words in the content
3. `n_unique_tokens`: Rate of unique words in the content
4. `n_non_stop_words`: Rate of non-stop words in the content
5. `n_non_stop_unique_tokens`: Rate of unique non-stop words in the content
6. `num_hrefs`: Number of links
7. `num_self_hrefs`: Number of links to other articles published by Mashable
8. `num_imgs`: Number of images

9. num\_videos: Number of videos
  10. average\_token\_length: Average length of the words in the content
  11. num\_keywords: Number of keywords in the metadata
  12. data\_channel\_is\_lifestyle: Is data channel 'Lifestyle'?
  13. data\_channel\_is\_elnetertainment: Is data channel 'Entertainment'?
  14. data\_channel\_is\_bus: Is data channel 'Business'?
  15. data\_channel\_is\_socmed: Is data channel 'Social Media'?
  16. data\_channel\_is\_tech: Is data channel 'Tech'?
  17. data\_channel\_is\_world: Is data channel 'World'?
  18. kw\_min\_min: Worst keyword (min. shares)
  19. kw\_max\_min: Worst keyword (max. shares)
  20. kw\_avg\_min: Worst keyword (avg. shares)
  21. kw\_min\_max: Best keyword (min. shares)
  22. kw\_max\_max: Best keyword (max. shares)
  23. kw\_avg\_max: Best keyword (avg. shares)
  24. kw\_min\_avg: Avg. keyword (min. shares)
  25. kw\_max\_avg: Avg. keyword (max. shares)
  26. kw\_avg\_avg: Avg. keyword (avg. shares)
  27. self\_reference\_min\_shares: Min. shares of referenced articles in Mashable
  28. self\_reference\_max\_shares: Max. shares of referenced articles in Mashable
  29. self\_reference\_avg\_shares: Avg. shares of referenced articles in Mashable
  30. weekday\_is\_monday: Was the article published on a Monday?
  31. weekday\_is\_tuesday: Was the article published on a Tuesday?
  32. weekday\_is\_wednesday: Was the article published on a Wednesday?
  33. weekday\_is\_thursday: Was the article published on a Thursday?
  34. weekday\_is\_friday: Was the article published on a Friday?
  35. weekday\_is\_saturday: Was the article published on a Saturday?
  36. weekday\_is\_sunday: Was the article published on a Sunday?
  37. is\_weekend: Was the article published on the weekend?
  38. LDA\_00: Closeness to LDA topic 0
  39. LDA\_01: Closeness to LDA topic 1
  40. LDA\_02: Closeness to LDA topic 2
  41. LDA\_03: Closeness to LDA topic 3
  42. LDA\_04: Closeness to LDA topic 4
  43. global\_subjectivity: Text subjectivity
  44. global\_sentiment\_polarity: Text sentiment polarity
  45. global\_rate\_positive\_words: Rate of positive words in the content
  46. global\_rate\_negative\_words: Rate of negative words in the content
  47. rate\_positive\_words: Rate of positive words among non-neutral tokens
  48. rate\_negative\_words: Rate of negative words among non-neutral tokens
  49. avg\_positive\_polarity: Avg. polarity of positive words
  50. min\_positive\_polarity: Min. polarity of positive words
  51. max\_positive\_polarity: Max. polarity of positive words
  52. avg\_negative\_polarity: Avg. polarity of negative words
  53. min\_negative\_polarity: Min. polarity of negative words
  54. max\_negative\_polarity: Max. polarity of negative words
  55. title\_subjectivity: Title subjectivity
  56. title\_sentiment\_polarity: Title polarity
  57. abs\_title\_subjectivity: Absolute subjectivity level
  58. abs\_title\_sentiment\_polarity: Absolute polarity level
-

Impute missing data-points with their mean. What is  $n$  and  $p$ ?

The data does not contain missing values, so imputation is not required.

$$n = 39644$$

$$p = 58$$

Standardize the numerical predictors using equation (6.6) in the ISLR book.

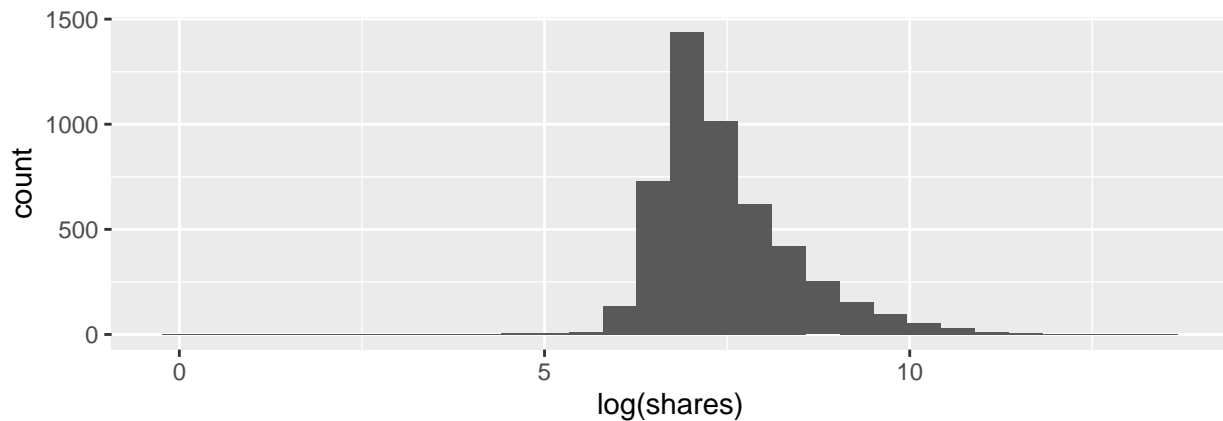

**Equation 6.6:**

$$\tilde{x}_{ij} = \frac{x_{ij}}{\sqrt{\frac{1}{n} \sum_{i=1}^n (x_{ij} - \bar{x}_j)^2}}$$

```
data <- data %>%  
  select(-shares) %>%  
  mutate_all(function(x) {x / sd(x)}) %>%  
  mutate(shares = data$shares)
```

**Distribution of  $y$**

```
hist_y <- data %>%  
  ggplot(aes(x=shares))+  
  geom_histogram(bins = 30)  
  
hist_log_y <- data %>%  
  ggplot(aes(x=log(shares)))+  
  geom_histogram(bins = 30)  
  
grid.arrange(hist_y, hist_log_y, nrow=2)
```



```
data <- data %>%
  mutate(shares=log(shares))

data %>%
  glimpse()
```

```
## Observations: 5,000
## Variables: 59
## $ n_tokens_title      <dbl> 3.757529, 5.166603, 4.696912, 3.75752...
## $ n_tokens_content    <dbl> 0.6308615, 0.3724882, 1.3198571, 0.78...
## $ n_unique_tokens     <dbl> 4.505415, 4.780077, 3.230670, 4.25455...
## $ n_non_stop_words    <dbl> 6.213983, 6.213983, 6.213983, 6.21398...
## $ n_non_stop_unique_tokens <dbl> 4.985482, 5.522474, 4.253439, 4.51426...
## $ num_hrefs           <dbl> 0.4520988, 0.5425185, 0.5425185, 0.27...
## $ num_self_hrefs      <dbl> 0.2619090, 0.7857269, 0.7857269, 0.78...
## $ num_imgs            <dbl> 0.1250814, 0.1250814, 0.1250814, 1.25...
## $ num_videos          <dbl> 0.239961, 0.000000, 0.000000, 0.00000...
## $ average_token_length <dbl> 6.468896, 5.413295, 5.417059, 6.30705...
## $ num_keywords        <dbl> 3.184113, 4.776170, 4.776170, 3.71479...
## $ data_channel_is_lifestyle <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...
## $ data_channel_is_entertainment <dbl> 0.000000, 0.000000, 0.000000, 2.59590...
## $ data_channel_is_bus <dbl> 2.771672, 2.771672, 0.000000, 0.00000...
## $ data_channel_is_socmed <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...
## $ data_channel_is_tech <dbl> 0.000000, 0.000000, 2.611734, 0.00000...
## $ data_channel_is_world <dbl> 0.000000, 0.000000, 0.000000, 0.00000, 2...
## $ kw_min_min          <dbl> -0.01450154, 3.14683521, 0.05800618, ...
## $ kw_max_min          <dbl> 0.10223094, 0.22377504, 0.07184492, 0...
```

```

## $ kw_avg_min <dbl> 0.24407628, 0.62642821, 0.16799824, 0...
## $ kw_min_max <dbl> 0.1424320, 0.0000000, 0.0000000, 0.29...
## $ kw_max_max <dbl> 3.963373, 2.904029, 3.963373, 3.96337...
## $ kw_avg_max <dbl> 2.8477590, 0.7746449, 1.3885462, 1.98...
## $ kw_min_avg <dbl> 1.4141508, 0.0000000, 0.0000000, 1.17...
## $ kw_max_avg <dbl> 0.5822269, 0.8806021, 0.6095905, 0.94...
## $ kw_avg_avg <dbl> 1.928369, 2.099917, 1.938905, 2.80282...
## $ self_reference_min_shares <dbl> 0.55256923, 0.17047349, 0.09993273, 0...
## $ self_reference_max_shares <dbl> 0.23231072, 0.14086927, 0.12356953, 0...
## $ self_reference_avg_shares <dbl> 0.43392180, 0.18157011, 0.12925330, 0...
## $ weekday_is_monday <dbl> 0.000000, 0.000000, 0.000000, 0.00000...
## $ weekday_is_tuesday <dbl> 0.00000, 2.59926, 0.00000, 0.00000, 0...
## $ weekday_is_wednesday <dbl> 2.544717, 0.000000, 0.000000, 2.54471...
## $ weekday_is_thursday <dbl> 0.000000, 0.000000, 2.576164, 0.00000...
## $ weekday_is_friday <dbl> 0.000000, 0.000000, 0.000000, 0.00000...
## $ weekday_is_saturday <dbl> 0.00000, 0.00000, 0.00000, 0.00000, 0...
## $ weekday_is_sunday <dbl> 0.000000, 0.000000, 0.000000, 0.00000...
## $ is_weekend <dbl> 0.000000, 0.000000, 0.000000, 0.00000...
## $ LDA_00 <dbl> 2.97415241, 0.96870559, 0.50493613, 0...
## $ LDA_01 <dbl> 0.14810054, 0.60831183, 0.09873790, 2...
## $ LDA_02 <dbl> 0.11703301, 0.07783261, 0.07782465, 0...
## $ LDA_03 <dbl> 0.11417523, 0.07611679, 0.07720218, 0...
## $ LDA_04 <dbl> 0.40764143, 1.97650377, 2.80658497, 0...
## $ global_subjectivity <dbl> 2.798987, 4.148770, 3.839035, 3.89829...
## $ global_sentiment_polarity <dbl> 2.3225839, 2.2398209, 2.5514693, 1.67...
## $ global_rate_positive_words <dbl> 2.1917583, 3.0371344, 2.7618820, 2.55...
## $ global_rate_negative_words <dbl> 0.3211055, 0.0000000, 0.1534811, 1.80...
## $ rate_positive_words <dbl> 4.887703, 5.332040, 5.154305, 3.70924...
## $ rate_negative_words <dbl> 0.5307563, 0.0000000, 0.2123025, 1.93...
## $ avg_positive_polarity <dbl> 3.382688, 2.892794, 3.987211, 3.75457...
## $ min_positive_polarity <dbl> 0.4613762, 0.6920643, 0.4613762, 0.46...
## $ max_positive_polarity <dbl> 2.057821, 3.292513, 3.498295, 3.29251...
## $ avg_negative_polarity <dbl> -0.9819048, 0.0000000, -1.5710477, -2...
## $ min_negative_polarity <dbl> -0.4286202, 0.0000000, -0.6857923, -1...
## $ max_negative_polarity <dbl> -1.3309854, 0.0000000, -2.1295767, -1...
## $ title_subjectivity <dbl> 0.0000000, 1.9620591, 1.3873145, 0.00...
## $ title_sentiment_polarity <dbl> 0.0000000, 0.8109844, 0.5160810, 0.00...
## $ abs_title_subjectivity <dbl> 2.6523409, 0.7578117, 0.2411219, 2.65...
## $ abs_title_sentiment_polarity <dbl> 0.0000000, 0.9503838, 0.6047897, 0.00...
## $ shares <dbl> 7.313220, 8.556414, 8.342840, 6.53087...

```

For each  $n_{train} = 0.8n$ , repeat the following 100 times, do the following for the different models mentioned below.

- Randomly split the dataset into two mutually exclusive datasets  $D_{test}$  and  $D_{train}$  with size  $n_{test}$  and  $n_{train}$  such that  $n_{train} + n_{test} = n$ .
- Use  $D_{train}$  to fit lasso, elastic-net  $\alpha = 0.5$ , ridge, and random forests.
- Tune the  $\lambda$ s using 10-fold CV.
- For each estimated model calculate

$$R_{test}^2 = 1 - \frac{\frac{1}{n_{test}} \sum_{i \in D_{test}} (y_i - \hat{y}_i)^2}{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2}$$

```

n <- 5000
p <- 58

sample_and_train <- function(plots=F) {
  # (a)
  n_train <- as.integer(0.8 * n)
  n_test = n - n_train
  train_inds <- sample.int(n, n_train)
  D_train <- data[train_inds, ]
  D_test <- data[-train_inds, ]
  # (b) & (c)
  X_train <- select(D_train, -shares) %>% data.matrix()
  X_test <- select(D_test, -shares) %>% data.matrix()
  y_train <- D_train$shares
  y_test <- D_test$shares
  y <- data$shares
  # (d)
  ## lasso
  lasso_time_start <- Sys.time()
  lasso_cv <- cv.glmnet(X_train, y_train, alpha = 1)
  lasso_fit <- glmnet(X_train, y_train, alpha = 1, lambda = lasso_cv$lambda.min)
  lasso_y_train_hat <- predict(lasso_fit, X_train)
  lasso_y_test_hat <- predict(lasso_fit, X_test)
  lasso_resid_train <- as.vector(y_train - lasso_y_train_hat)
  lasso_resid_test <- as.vector(y_test - lasso_y_test_hat)
  lasso_Rsq_train <- 1 - mean((lasso_resid_train)^2) / mean((y - mean(y))^2)
  lasso_Rsq_test <- 1 - mean((lasso_resid_test)^2) / mean((y - mean(y))^2)
  lasso_time <- Sys.time() - lasso_time_start

  ## ridge
  ridge_time_start <- Sys.time()
  ridge_cv <- cv.glmnet(X_train, y_train, alpha = 0)
  ridge_fit <- glmnet(X_train, y_train, alpha = 0, lambda = ridge_cv$lambda.min)
  ridge_y_train_hat <- predict(ridge_fit, X_train)
  ridge_y_test_hat <- predict(ridge_fit, X_test)
  ridge_resid_train <- as.vector(y_train - ridge_y_train_hat)
  ridge_resid_test <- as.vector(y_test - ridge_y_test_hat)
  ridge_Rsq_train <- 1 - mean((ridge_resid_train)^2) / mean((y - mean(y))^2)
  ridge_Rsq_test <- 1 - mean((ridge_resid_test)^2) / mean((y - mean(y))^2)
  ridge_time <- Sys.time() - ridge_time_start

  ## elastic-net
  elnet_time_start <- Sys.time()
  elnet_cv <- cv.glmnet(X_train, y_train, alpha = 0.5)
  elnet_fit <- glmnet(X_train, y_train, alpha = 0.5, lambda = elnet_cv$lambda.min)
  elnet_y_train_hat <- predict(elnet_fit, X_train)
  elnet_y_test_hat <- predict(elnet_fit, X_test)
  elnet_resid_train <- as.vector(y_train - elnet_y_train_hat)
  elnet_resid_test <- as.vector(y_test - elnet_y_test_hat)
  elnet_Rsq_train <- 1 - mean((elnet_resid_train)^2) / mean((y - mean(y))^2)
  elnet_Rsq_test <- 1 - mean((elnet_resid_test)^2) / mean((y - mean(y))^2)
  elnet_time <- Sys.time() - elnet_time_start
}

```

```

## random forest
rf_time_start <- Sys.time()
rf_fit <- randomForest(X_train, y_train, mtry = sqrt(p), importance = T)
rf_y_train_hat <- predict(rf_fit, X_train)
rf_y_test_hat <- predict(rf_fit, X_test)
rf_resid_train <- y_train - rf_y_train_hat
rf_resid_test <- y_test - rf_y_test_hat
rf_Rsq_train <- 1 - mean((rf_resid_train)^2) / mean((y - mean(y))^2)
rf_Rsq_test <- 1 - mean((rf_resid_test)^2) / mean((y - mean(y))^2)
rf_time <- Sys.time() - rf_time_start

if(plots) {
  ## 10 fold CV Plots
  plot(lasso_cv, sub = paste("Lasso:", lasso_cv$lambda.min))
  plot(ridge_cv, sub = paste("Ridge", ridge_cv$lambda.min))
  plot(elnet_cv, sub = paste("Elastic Net:", elnet_cv$lambda.min))

  ## Residuals Boxplots
  resid_train <- data.frame(lasso = lasso_resid_train, ridge = ridge_resid_train,
                           elnet = elnet_resid_train, rf = rf_resid_train, dataset="train")
  resid_test <- data.frame(lasso = lasso_resid_test, ridge = ridge_resid_test,
                           elnet = elnet_resid_test, rf = rf_resid_test, dataset="test")
  resid_models <- rbind(resid_train, resid_test)
  resid_plot <- resid_models %>%
    gather(model, residuals, lasso:rf) %>%
    ggplot(aes(x=model, y=residuals, fill=model)) +
    geom_boxplot() +
    facet_wrap(~dataset)

  # resid_plot < ggplotly(resid_plot) # Un-comment for ggplotly plot
  print(resid_plot)
}


Rsq_train <- list(lasso = lasso_Rsq_train, ridge = ridge_Rsq_train,
                 elnet = elnet_Rsq_train, rf = rf_Rsq_train)
Rsq_test <- list(lasso = lasso_Rsq_test, ridge = ridge_Rsq_test,
                elnet = elnet_Rsq_test, rf = rf_Rsq_test)
times <- list(lasso = lasso_time, ridge = ridge_time,
              elnet = elnet_time, rf = rf_time)

return(list(Rsq_train, Rsq_test, times))
}


set.seed(2)

one_sample <- sample_and_train(plots = T)

```




Lasso: 0.0042965206503351




Ridge 0.12524585722488





Elastic Net: 0.00859304130067019



```
set.seed(2)
M <- 100
Rsqr_time_models <- replicate(M, unlist(sample_and_train(plots = F))) %>% t()
```

```


Rsqr_train <- Rsqr_time_models[, 1:4] %>% data.frame() %>% mutate(dataset="train")
Rsqr_test <- Rsqr_time_models[, 5:8] %>% data.frame() %>% mutate(dataset="test")
time_models <- Rsqr_time_models[, 9:12] %>% data.frame()

Rsqr_plot_data <- rbind(Rsqr_train, Rsqr_test) %>%
  gather(model, Rsqr, lasso:rf)

Rsqr_plot <- Rsqr_plot_data %>%
  mutate(dataset = factor(dataset, levels=c("train", "test"))) %>%
  ggplot(aes(x = model, y = Rsqr)) +
  geom_boxplot(aes(fill = model)) +
  facet_wrap(~ dataset, scales = "free_y") +
  labs(x = "Model", y = "R-Squared")

#Rsqr_plot <- ggplotly(Rsqr_plot) # Un-comment for ggplotly plot
Rsqr_plot

```



```

sample_and_train_bs <- function() {
  bs_inds <- sample(n, replace=T)

  X_bs <- data[bs_inds, ] %>%
    select(-shares) %>%
    data.matrix()
  y_bs <- data$shares[bs_inds]

  # bootstrap lasso

```

```

lasso_cv <- cv.glmnet(X_bs, y_bs, alpha = 1)
lasso_fit <- glmnet(X_bs, y_bs, alpha = 1, lambda = lasso_cv$lambda.min)
beta_lasso_bs <- as.vector(lasso_fit$beta)
# bootstrap ridge
ridge_cv <- cv.glmnet(X_bs, y_bs, alpha = 0)
ridge_fit <- glmnet(X_bs, y_bs, alpha = 0, lambda = ridge_cv$lambda.min)
beta_ridge_bs <- as.vector(ridge_fit$beta)
# bootstrap elastic-net
elnet_cv <- cv.glmnet(X_bs, y_bs, alpha = 0.5)
elnet_fit <- glmnet(X_bs, y_bs, alpha = 0.5, lambda = elnet_cv$lambda.min)
beta_elnet_bs <- as.vector(elnet_fit$beta)
# bootstrap random forest
rf_fit <- randomForest(X_bs, y_bs, mtry = sqrt(p), importance = TRUE)
beta_rf_bs <- as.vector(rf_fit$importance[,1])

return(list(beta_lasso_bs=beta_lasso_bs, beta_ridge_bs=beta_ridge_bs,
            beta_elnet_bs=beta_elnet_bs, beta_rf_bs=beta_rf_bs))
}
set.seed(3)
bootstrapSamples = 100
beta_models <- replicate(bootstrapSamples, sample_and_train_bs())

beta_lasso_bs <- do.call(rbind, beta_models[1,])
beta_ridge_bs <- do.call(rbind, beta_models[2,])
beta_elnet_bs <- do.call(rbind, beta_models[3,])
beta_rf_bs <- do.call(rbind, beta_models[4,])

# calculate bootstrapped standard errors
lasso_bs_sd <- apply(beta_lasso_bs, 2, "sd")
ridge_bs_sd <- apply(beta_ridge_bs, 2, "sd")
rf_bs_sd <- apply(beta_rf_bs, 2, "sd")
elnet_bs_sd <- apply(beta_elnet_bs, 2, "sd")

X <- data %>% select(-shares) %>% data.matrix()
y <- data$shares

# fit lasso to the whole data
lasso_cv <- cv.glmnet(X, y, alpha = 1)
lasso_fit <- glmnet(X, y, alpha = 1, lambda = lasso_cv$lambda.min)
# fit ridge to the whole data
ridge_cv <- cv.glmnet(X, y, alpha = 0)
ridge_fit <- glmnet(X, y, alpha = 0, lambda = ridge_cv$lambda.min)
# fit elnet to the whole data
elnet_cv <- cv.glmnet(X, y, alpha = 0.5)
elnet_fit <- glmnet(X, y, alpha = 0.5, lambda = elnet_cv$lambda.min)
# fit rf to the whole data
rf_fit <- randomForest(X, y, ntree = 1, mtry = sqrt(p), importance = TRUE)

features <- 1:length(names(X[1,])) %>% as.factor()

betaS_lasso <- data.frame(feature = features, value = as.vector(lasso_fit$beta),
                          error = 2*lasso_bs_sd, model = "lasso")
betaS_ridge <- data.frame(feature = features, value = as.vector(ridge_fit$beta),


```

```

      error = 2*ridge_bs_sd, model = "ridge")
betaS_elnet <- data.frame(feature = features, value = as.vector(elnet_fit$beta),
      error = 2*elnet_bs_sd, model = "elnet")
betaS_rf <- data.frame(feature = features, value = as.vector(rf_fit$importance[,1]),
      error = 2*rf_bs_sd, model = "rf")

betaS_models <- rbind(betaS_lasso, betaS_ridge, betaS_elnet, betaS_rf)
betaS_plot <- betaS_models %>%
  ggplot(aes(x=feature, y=value)) +
  geom_bar(stat = "identity", fill="white", colour="black") +
  geom_errorbar(aes(ymin=value-error, ymax=value+error), width=.2) +
  facet_wrap(~ model, scales = "free_y", ncol = 1)
#betaS_plot <- ggplotly(betaS_plot) # Un-comment for ggplotly plot
betaS_plot

```




```

time_plot <- time_models %>%
  gather(model, time, lasso:rf) %>%
  ggplot(aes(x=model, y=time, fill=model)) +
  geom_boxplot()

#time_plot <- ggplotly(time_plot) # Un-comment for ggplotly plot
time_plot

```



```
avg_model_times <- time_models %>%
  gather(model, time, lasso:rf) %>%
  group_by(model) %>%
  summarize(mean=mean(time), median=median(time), min=min(time), max=max(time))
kable(avg_model_times, digits=2)
```

model	mean	median	min	max
elnet	0.29	0.28	0.24	0.46
lasso	0.30	0.29	0.24	0.46
rf	27.96	27.31	26.85	32.94
ridge	0.30	0.30	0.26	0.47